

# **Subject:** Capstone Machine Learning Deployment Assignment

## **Project :** Car Price Prediction API with FastAPI & Google Cloud Run

---

### **1. Executive Summary**

This project aims to build and deploy a machine learning model to predict car prices based on customer demographics and vehicle specifications. The solution uses Python's FastAPI framework for serving predictions and is deployed using Google Cloud Run for scalability and availability. The model is trained on a car sales dataset, preprocessed for quality, and exposed via a REST API.

---

### **2. Problem Statement & Dataset Description**

**Problem:** Predict the price of a car based on inputs like age, annual income, engine type, company, etc.

**Dataset Source:** Kaggle - [Car Sales Dataset](#)

#### **Features Used:**

- Age
- Annual Income
- Company
- Model
- Engine
- Transmission
- Dealer Region

#### **Target Variable:**

- Price (\$)

	Car_id	Date	Customer Name	Gender	Annual Income	Dealer Name	Company	Model	Engine	Transmission	Color	Price (\$)	Dealer_No	Body Style	Phone
0	C_CND_000001	01-02-2022	Geraldine	Male	13500	Buddy Storbeck's Diesel Service Inc	Ford	Expedition	DoubleA Overhead Camshaft	Auto	Black	26000	06457-3834	SUV	8264678
1	C_CND_000002	01-02-2022	Gia	Male	1480000	C & M Motors Inc	Dodge	Durango	DoubleA Overhead Camshaft	Auto	Black	19000	60504-7114	SUV	6848189
2	C_CND_000003	01-02-2022	Gianna	Male	1035000	Capitol KIA	Cadillac	Eldorado	Overhead Camshaft	Manual	Red	31500	38701-8047	Passenger	7298798
3	C_CND_000004	01-02-2022	Giselle	Male	13500	Chrysler of Tri-Cities	Toyota	Celica	Overhead Camshaft	Manual	Pale White	14000	99301-3882	SUV	6257557
4	C_CND_000005	01-02-2022	Grace	Male	1465000	Chrysler Plymouth	Acura	TL	DoubleA Overhead Camshaft	Auto	Red	24500	53546-9427	Hatchback	7081483

### 3. Model Development Process

#### Preprocessing:

- Missing value imputation
- Label encoding for categorical variables
- Feature scaling with StandardScaler

#### Model:

- RandomForestRegressor (with GridSearchCV)

#### Evaluation Metrics:

- $R^2$  Score
- Mean Absolute Error (MAE)

#### Performance:



```

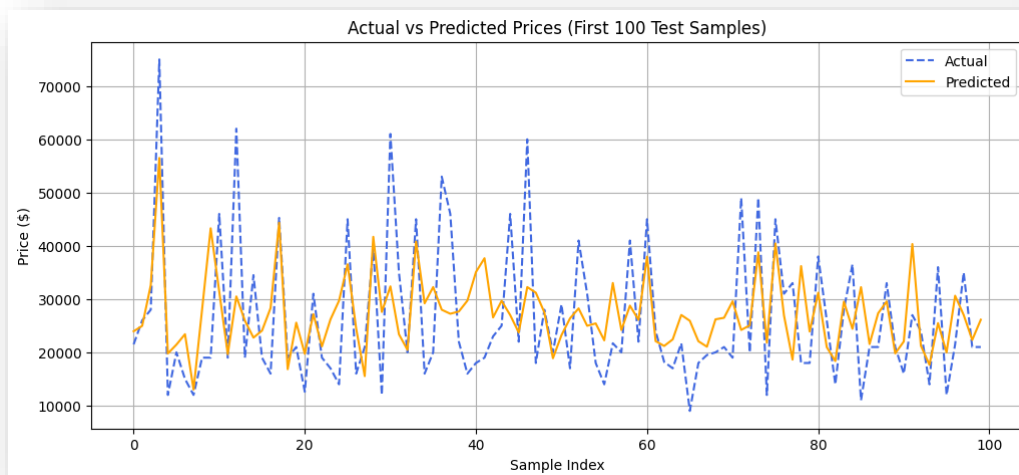

Training Set Evaluation
MAE: 8990.69
RMSE: 11876.50
R²: 0.2705



Validation Set Evaluation
MAE: 9253.97
RMSE: 12242.53
R²: 0.2264

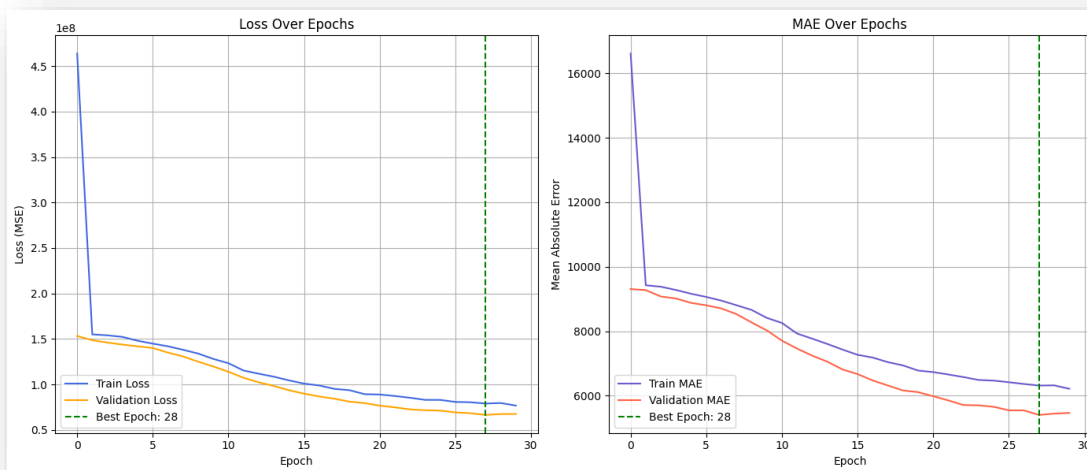
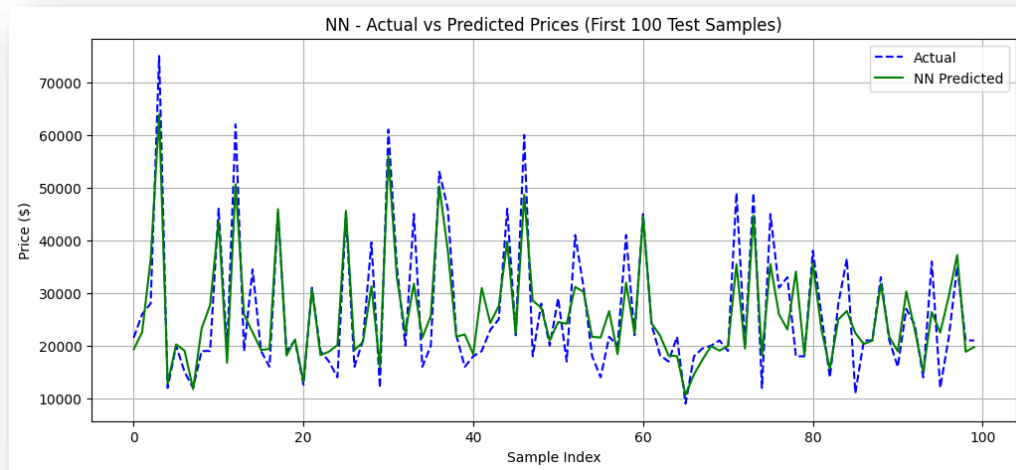
Test Set Evaluation
MAE: 9047.68
RMSE: 11872.09
R²: 0.2248

```









## 4. Deployment Architecture

### Tools & Services:

- FastAPI
- Docker
- Google Cloud Build
- Google Cloud Run
- Cloud Container Registry

### Deployment Steps:

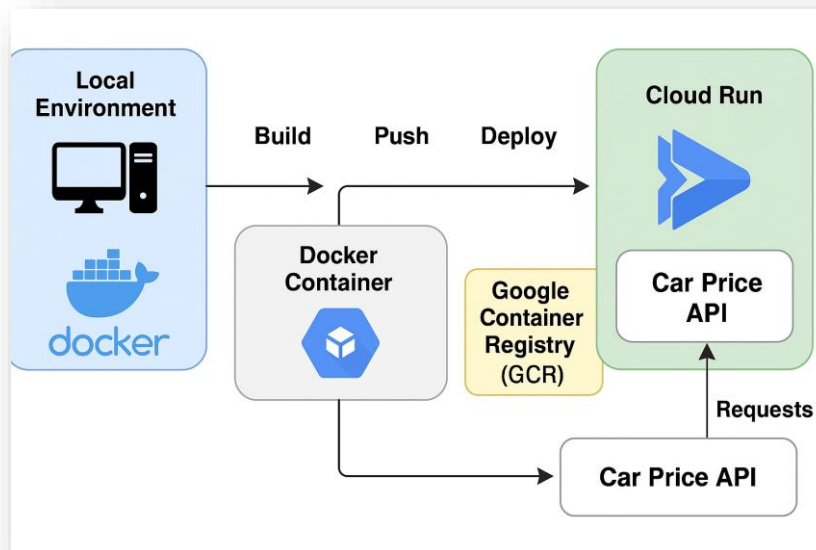
1. Train and pickle model

2. Create FastAPI app in main.py
3. Write Dockerfile
4. Build and push image to GCR
5. Deploy to Cloud Run with memory set to 1Gi

**Live API Endpoint:**

<https://car-price-api-481594874299.us-central1.run.app/docs>

Diagram showing architecture (local -> Docker -> GCR -> Cloud Run)



Screenshot of successful deployment confirmation in terminal

```
A:\Seneca\Sem 1\Machine Learning\PROJECT 3\cp_individual>gcloud run deploy car-price-api --image gcr.io/carprice-capstone-ind/car-price-api --platform managed --region us-central1 --allow-unauthenticated --memory 1Gi
Deploying container to Cloud Run service [car-price-api] in project [carprice-capstone-ind] region [us-central1]
OK Deploying... Done.
OK Creating Revision...
OK Routing traffic...
OK Setting IAM Policy...
Done.
Service [car-price-api] revision [car-price-api-00005-45q] has been deployed and is serving 100 percent of traffic.
Service URL: https://car-price-api-481594874299.us-central1.run.app
A:\Seneca\Sem 1\Machine Learning\PROJECT 3\cp_individual>
```

---

## 6. Challenges & Solutions

**Challenge:** Memory limit exceeded 512MiB

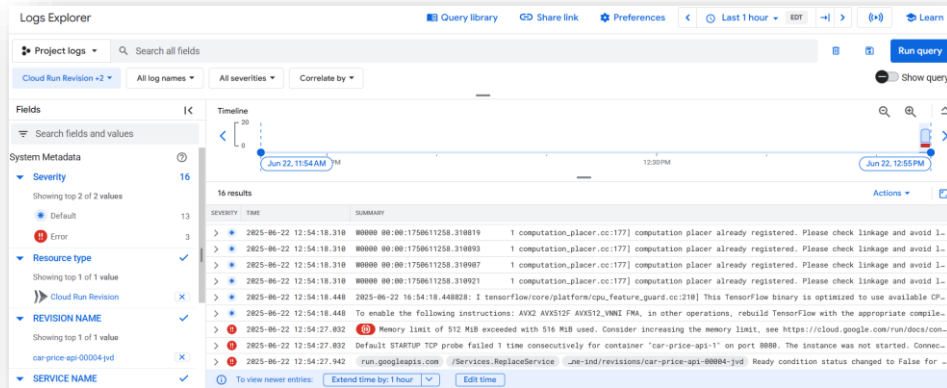
- **Solution:** Increased to `--memory 1Gi` in deploy command

**Challenge:** App not listening on port 8080

- **Solution:** Ensured FastAPI uses port=8080 in CMD and EXPOSE in Dockerfile

**Challenge:** Timeout errors

- **Solution:** Verified logs and updated startupProbe configuration if needed



Error logs or deployment failure messages from Cloud Console

## 7. Conclusion & Future Work

This project demonstrates the ability to train, serve, and deploy an ML model using modern cloud tools. The app is now scalable and available publicly via REST API. Future enhancements include:

- Adding authentication
- UI dashboard for inputs
- Training on additional features like fuel type, mileage, etc.



Working screenshot of API on Swagger UI or Postman.

